

## Optimization:

To do gradient descent, PyTorch provides optimization packages. This package contains standard optimizers.

PyTorch has a `torch.optim` package that is used to implement various optimization algorithms. Most common optimizers used today like SGD, Adam, ReLu etc are already available within this package. It also has the robustness to add more sophisticated optimizers as per your need.

In order to use `torch.optim` we will have to construct an optimizer object which will store the current state and will update the parameters of the optimizers based on the computed gradients.

In order to initialize an optimizer this is the syntax to be followed -

CODE IN GOOGLE COLAB

## Optimization algorithms that are already available in the `torch.optim` package are:

Adadelta - <https://arxiv.org/abs/1212.5701>

Adagrad - <http://jmlr.org/papers/v12/duchi11a.html>

Adam - <https://arxiv.org/abs/1412.6980>

Adamax - <https://arxiv.org/abs/1412.6980>

ASGD - <http://dl.acm.org/citation.cfm?id=131098>

LBFGS -

RMSprop - Proposed by G. Hinton in his [course](#).

The centered version first appears in [Generating Sequences With Recurrent Neural Networks](#).

SGD - <http://www.cs.toronto.edu/~hinton/absps/momentum.pdf>

These algorithms can be used with just one line of code. The syntax for using these algorithms and their parameter definition can be found in PyTorch official documentation here -

<https://pytorch.org/docs/stable/optim.html#algorithms>

### **Adjusting learning\_rate:**

What is learning rate and why it needs to be adjusted?

PyTorch has a very good documentation explaining how to adjust the learning rate in PyTorch. The documentation can be found here -

<https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate>